

Examining nonlinear methods for toxicity prediction of carboxylic acids

Abdelsalam A

The findings of *Improved QSAR Analysis of the Toxicity of Aliphatic Carboxylic Acids* found improvements in predicting aquatic toxicity of carboxylic acids by using a particular set of molecular and topological parameters. The statistical methods employed were linear, and polynomial in nature which relied upon removing predictors. We attempt to establish a stronger predictive relationship using purely nonlinear methods while retaining all predictors.

The dataset can be found at: *Toxicity of Carboxylic Acids Data*

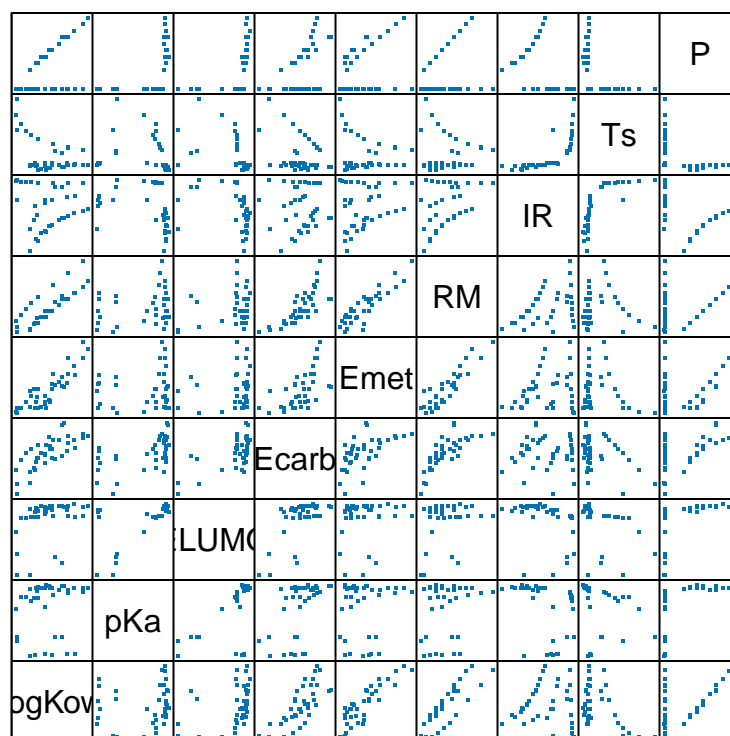
Firstly, we take a peak at the dataset:

```
head(toxicity)
```

| ## | toxicity | logKow | pKa | ELUMO | Ecarb | Emet | RM | IR | Ts | P |
|------|----------|--------|------|-------|---------|--------|-------|-------|------|-------|
| ## 1 | -0.15 | 1.68 | 1.00 | 4.81 | 17.8635 | 1.4838 | 31.36 | 1.425 | 31.3 | 12.43 |
| ## 2 | -0.33 | 0.94 | 0.98 | 4.68 | 16.9491 | 0.0000 | 22.10 | 1.408 | 30.4 | 8.76 |
| ## 3 | -0.34 | 1.16 | 0.96 | 4.86 | 17.1806 | 0.2778 | 26.73 | 1.418 | 30.9 | 10.59 |
| ## 4 | 0.03 | 2.75 | 1.00 | 4.83 | 18.4794 | 3.5836 | 40.63 | 1.435 | 31.8 | 16.10 |
| ## 5 | -0.57 | 0.79 | 0.97 | 4.80 | 16.8022 | 1.0232 | 22.14 | 1.411 | 32.5 | 8.77 |
| ## 6 | 0.08 | 2.64 | 1.01 | 4.90 | 18.3937 | 3.7145 | 40.63 | 1.435 | 31.8 | 16.10 |


We want to examine the pairwise relationships of features, so we pre-process via centering, scaling, followed by plotting the scatter-plot matrix, and correlation matrix

```
# Centering & Scaling
scaledToxicity <- preProcess(toxicity[, -1], method = c("center", "scale"))
csData <- predict(scaledToxicity, newdata = toxicity[, -1])
# Scatterplot Matrix
splom(~csData, pch = 15, cex = .25, pscales = 0)
```



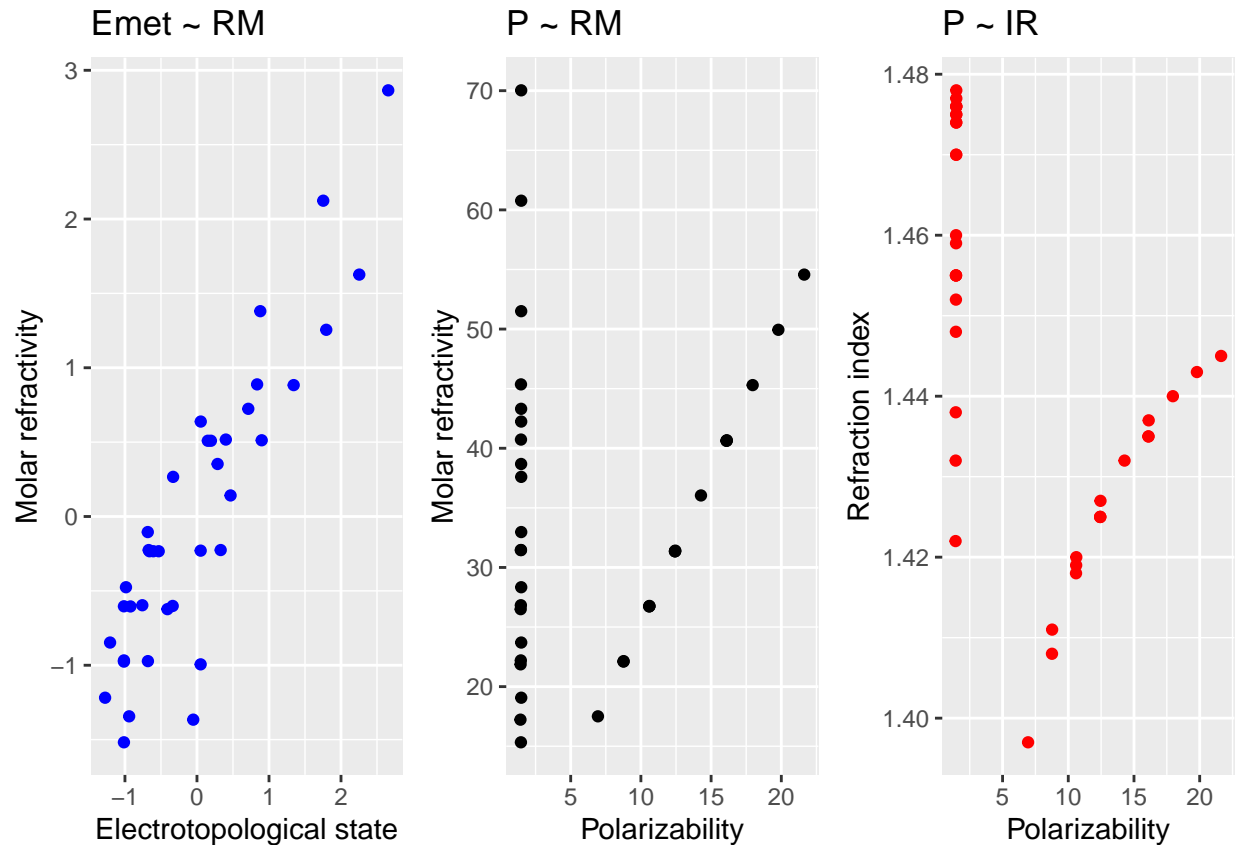
Scatter Plot Matrix

```
# Correlation Matrix
corrplot(cor(csData))
```



We observe a few select heavily-correlated non-linear features, such as **RM ~ Emet**, some with a mix of both linear and uncorrelated points, like **P ~ RM**, and some with non-linear and uncorrelated points, like **P ~ IR**

```
csData <- as.data.frame(csData)
plt1 <- ggplot(csData, aes(x=Emet, y = RM)) + geom_point(color = "blue") + labs(
  title = "Emet ~ RM",
  x = "Electrotopological state", y = "Molar refractivity",
)
plt2 <- ggplot(toxicity, aes(x=P, y = RM)) + geom_point() + labs(
  title = "P ~ RM",
  x = "Polarizability", y = "Molar refractivity",
)
plt3 <- ggplot(toxicity, aes(x=P, y = IR)) + geom_point(color = "red") + labs(
  title = "P ~ IR",
  x = "Polarizability", y = "Refraction index",
)
ggarrange(plt1, plt2, plt3, ncol = 3)
```



We see that some feature-feature relationships occur only after a certain threshold. For example, as polarizability increases after ~ 6.5 , molar refractivity increases with it linearly. A model that captures these sudden pattern changes is MARS.

Multivariate Adaptive Regression Splines (MARS)

```
set.seed(1)
# Make training and test sets
trainingRows <- createDataPartition(toxicity[,1], p = 0.7, list = FALSE)
# Training set
trainSet <- toxicity[trainingRows,]
# Testing set
testSet <- toxicity[-trainingRows,]
# 1-3 degree interactions
hyper_grid <- expand.grid(
  degree = 1:3,
  nprune = seq(2, 100, length.out = 10) %>% floor()
)


# MARS with 5-fold CV
marsModel <- train(
  toxicity ~ .,
  data = trainSet,
  method = "earth",
  metric = "RMSE",
  preProcess = c("BoxCox", "center", "scale"),
```

```

trControl = trainControl(method = "cv", number = 5),
tuneGrid = hyper_grid
)
# Plotting results
predictedValuesMars <- predict(marsModel, testSet)
predictedDataMars <- cbind(testSet[,1], predictedValuesMars)
colnames(predictedDataMars) <- c("expectedY", "predictedY")
ggplot(predictedDataMars, aes(x = expectedY, y = predictedY)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE)

## `geom_smooth()` using formula = 'y ~ x'

```



```

# R2 value
caret::R2(predictedValuesMars, testSet[,1])

##           y
## [1,] 0.7462671

```

Unfortunately, MARS didn't perform particularly well with an R^2 of 0.7463.

We follow with SVM; in particular, we try a radial basis, and linear kernel.

```

set.seed(1)
# Radial Basis kernel
svmRadialModel <- train(
  toxicity ~ .,
  data = trainSet,

```

```

method = "svmRadial",
preProc = c("BoxCox", "center", "scale"),
tuneLength = 14,
trControl = trainControl(method = "cv"))
# Linear kernel
svmLinearModel <- train(
  toxicity ~ .,
  data = trainSet,
  method = "svmLinear",
  preProc = c("BoxCox", "center", "scale"),
  tuneLength = 14,
  trControl = trainControl(method = "cv"))

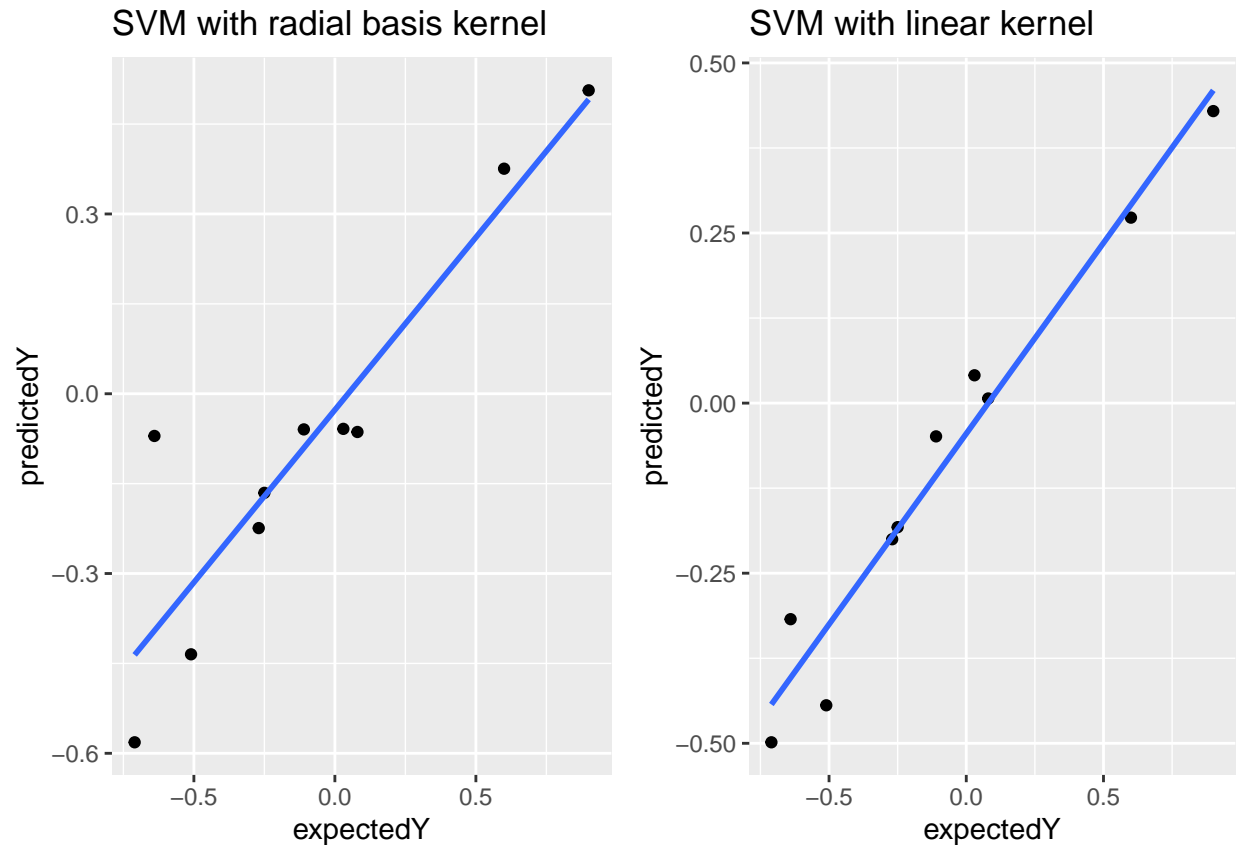
# Plotting results
predictedValuesRadial <- predict(svmRadialModel, testSet)
predictedDataRadial <- cbind(testSet[,1], predictedValuesRadial)
colnames(predictedDataRadial) <- c("expectedY", "predictedY")
pltSVMRadial <- ggplot(predictedDataRadial, aes(x = expectedY, y = predictedY)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  labs(title = "SVM with radial basis kernel")

predictedValuesLinear <- predict(svmLinearModel, testSet)
predictedDataLinear <- cbind(testSet[,1], predictedValuesLinear)
colnames(predictedDataLinear) <- c("expectedY", "predictedY")
pltSVMLinear <- ggplot(predictedDataLinear, aes(x = expectedY, y = predictedY)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  labs(title = "SVM with linear kernel")

ggarrange(pltSVMRadial, pltSVMLinear, ncol = 2)

## `geom_smooth()` using formula = 'y ~ x'
## `geom_smooth()` using formula = 'y ~ x'

```



```
# R2 value of SVM Radial
caret::R2(predictedValuesRadial, testSet[,1])
```

```
## [1] 0.8377675
```

```
# R2 value of SVM Linear
caret::R2(predictedValuesLinear, testSet[,1])
```

```
## [1] 0.9586954
```

The radial basis kernel performed mediocre at $R^2 = 0.8377675$, while the linear kernel performed fairly well at $R^2 = 0.9586954$ which beats all but two polynomial models in the original study.

Finally, we attempt a neural network, and K-Nearest-neighbours.

For NN, we can't have our correlated features crippling the model, so we apply PCA for pre-processing.

```
set.seed(1)
# Neural Network with PCA
nnetGrid <- expand.grid(.decay = c(0, 0.01, .1),
  .size = c(1:10),
  .bag = FALSE)
nnetModel <- train(toxicity ~ .,
  data = trainSet,
  method = "nnet",
  tuningGrid = nnetGrid,
  trControl = trainControl(method = "cv"),
  preProc = c("center", "scale", "pca"),
  linout = TRUE,
```

```

trace = FALSE,
MaxNWts = 10 * (ncol(toxicity[, -1]) + 1) + 10 + 1,
maxit = 500)

#K-Nearest-Neighbors

knnModel <- train(toxicity ~ .,
  data = trainSet,
  method = "knn",
  preProc = c("center", "scale"),
  tuneGrid = data.frame(.k = 1:20),
  trControl = trainControl(method = "cv"))

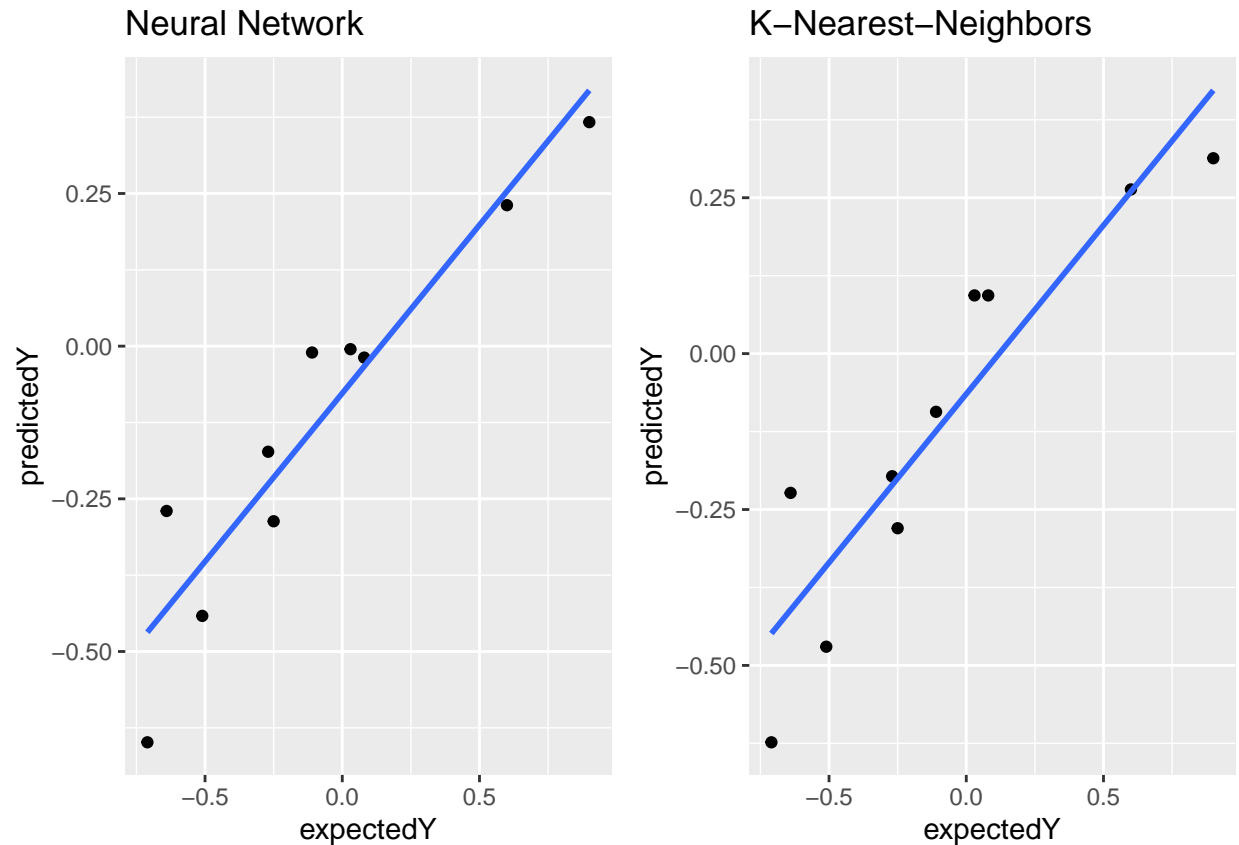
# Generating predictions - NN
predictedValuesNNNet <- predict(nnetModel, testSet)
predictedDataNNNet <- cbind(testSet[,1], predictedValuesNNNet)
colnames(predictedDataNNNet) <- c("expectedY", "predictedY")
pltNNNet <- ggplot(predictedDataNNNet, aes(x = expectedY, y = predictedY)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  labs(title = "Neural Network")

# Generating predictions - KNN
predictedValuesKNN <- predict(knnModel, testSet)
predictedDataKNN <- cbind(testSet[,1], predictedValuesKNN)
colnames(predictedDataKNN) <- c("expectedY", "predictedY")
pltKNN <- ggplot(predictedDataKNN, aes(x = expectedY, y = predictedY)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  labs(title = "K-Nearest-Neighbors")

# Plotting results
ggarrange(pltNNNet, pltKNN, ncol = 2)

## `geom_smooth()` using formula = 'y ~ x'
## `geom_smooth()` using formula = 'y ~ x'

```

```
# R2 value of NNet
caret::R2(predictedValuesNNet, testSet[,1])
```

```
## [1] 0.8847985
```

```
# R2 value of KNN
caret::R2(predictedValuesKNN, testSet[,1])
```

```
## [1] 0.8399075
```

We observe that both NNet, and KNN perform fairly mediocre with R^2 values of 0.8848 and 0.8399.

We conclude that the nonlinear methods examined in this analysis do not perform well enough to compete with the simpler, and better-performing linear methods described in the original study.